

## The Contents in SSD6 cover:

- 4.1 Measurement and Profiling
  Where to start? Where to stop? How-to?
- 4.2 Hot Spots

  Find the weak-knees!
- 4.3 Practical Hints

Join a tour...







# **Key Points:**

Trade-offs

- Making program run iast.
- Making program run correct!
- Which is more important?
- Computers double in speed every 18 months







#### **General Hints:**

- Understand what makes the program slow.
  - A few places account for a lot of time.
- Design and implement a solution
  - The most speedup for the least effort
- Iterate steps 1 and 2
  - ■until performance is satisfactory or,
  - until the remaining optimizations are not worth the effort







#### Where to measure:

- Where programs spend most of their time.
  - Don't optimize a under-dev. program.
- **CPU** time
  - For instruction execution
  - System time vs. User time
- Wall time
  - For all elapsed time
- Difference between 2&3







## **Measurement and Profiling**

### Time scale:

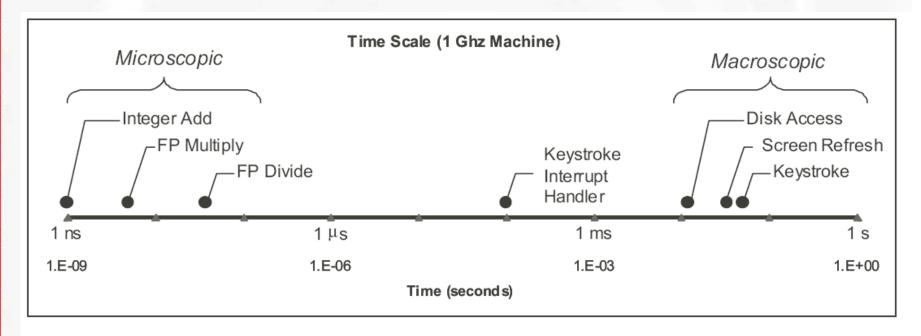


Figure 9.1: **Time Scale of Computer System Events.** The processor hardware works at a microscopic a time scale in which events having durations on the order of a few nanoseconds (ns). The OS must deal on a macroscopic time scale with events having durations on the order of a few milliseconds (ms).







### Time scale:

- Computers operate on two fundamentally different time scales
- At a microscopic level
  - To execute instructions at a rate of one or more per clock cycle, where each clock cycle requires only around one nanosecond (abbreviated "ns")
- □ On a macroscopic scale
  - processor must respond to external events that occur on time scales measured in milliseconds (abbreviated "ms")







## **Measurement and Profiling**

# For example:

- □ During video playback, the graphics display for most computers must be refreshed every 33 ms
- A world-record typist can only type keystrokes at a rate of around one every 50 milliseconds
- Disks typically require around 10 ms to initiate a disk transfer
- ☐ The processor continually switches between these many tasks on a macroscopic time scale, devoting around 5 to 20 *milliseconds* to each task at a time
- People cannot discern time durations shorter than around 100 *ms*. Within that time the processor can execute millions of instructions









- Computers have an external timer that periodically generates an interrupt signal to the processor
- The spacing between these interrupt signals is called the interval time
- When a timer interrupt occurs, the OS scheduler can choose to either resume the currently executing process or to switch to a different process.
- ☐ This interval must be set short enough to ensure that the processor will switch between tasks
- Typical timer intervals range between 1 and 10 milliseconds

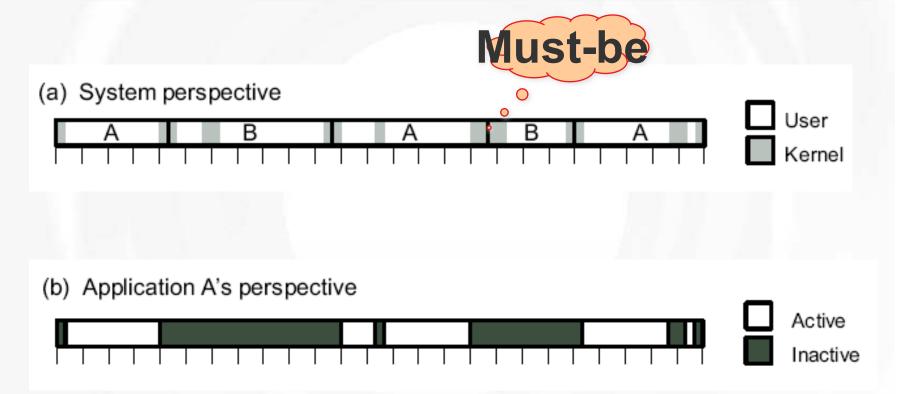






## **Measurement and Profiling**

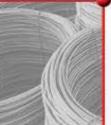
## The Flow of Time on a Computer System:











#### **Time from an Application Program's Perspective**

#### Active vs. Inactive

| A0 | Time | 0        | (0.00  | ms), | Duration | 3726508 | (6.776448 | ms) |
|----|------|----------|--------|------|----------|---------|-----------|-----|
| ΙO | Time | 3726508  | (6.78  | ms), | Duration | 275025  | (0.500118 | ms) |
| A1 | Time | 4001533  | (7.28  | ms), | Duration | 0       | (0.000000 | ms) |
| I1 | Time | 4001533  | (7.28  | ms), | Duration | 7598    | (0.013817 | ms) |
| A2 | Time | 4009131  | (7.29  | ms), | Duration | 5189247 | (9.436358 | ms) |
| I2 | Time | 9198378  | (16.73 | ms), | Duration | 251609  | (0.457537 | ms) |
| A3 | Time | 9449987  | (17.18 | ms), | Duration | 2250102 | (4.091686 | ms) |
| I3 | Time | 11700089 | (21.28 | ms), | Duration | 14116   | (0.025669 | ms) |
| A4 | Time | 11714205 | (21.30 | ms), | Duration | 2955974 | (5.375275 | ms) |
| I4 | Time | 14670179 | (26.68 | ms), | Duration | 248500  | (0.451883 | ms) |
| A5 | Time | 14918679 | (27.13 | ms), | Duration | 5223342 | (9.498358 | ms) |
| I5 | Time | 20142021 | (36.63 | ms), | Duration | 247113  | (0.449361 | ms) |
| A6 | Time | 20389134 | (37.08 | ms), | Duration | 5224777 | (9.500967 | ms) |
| I6 | Time | 25613911 | (46.58 | ms), | Duration | 254340  | (0.462503 | ms) |
| A7 | Time | 25868251 | (47.04 | ms), | Duration | 3678102 | (6.688425 | ms) |
| I7 | Time | 29546353 | (53.73 | ms), | Duration | 8139    | (0.014800 | ms) |
| A8 | Time | 29554492 | (53.74 | ms), | Duration | 1531187 | (2.784379 | ms) |
| 18 | Time | 31085679 | (56.53 | ms), | Duration | 248360  | (0.451629 | ms) |
| A9 | Time | 31334039 | (56.98 | ms), | Duration | 5223581 | (9.498792 | ms) |
| Ι9 | Time | 36557620 | (66.48 | ms), | Duration | 247395  | (0.449874 | ms) |
|    |      |          |        |      |          |         |           |     |





#### Time from an Application Program's Perspective

Active vs. Inactive

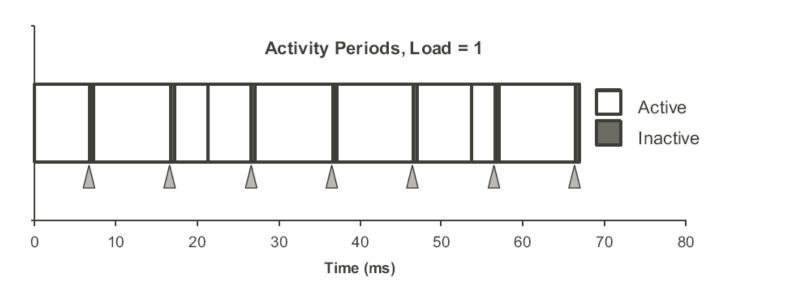


Figure 9.4: **Graphical Representation of Trace in Figure 9.3.** Timer interrupts are indicated with gray triangles.









- The operating system also uses the timer to record the cumulative time used by each process.
- This information provides a somewhat imprecise measure of program execution time
- A time segment refers to the period during which just one process executes.







## **Measurement and Profiling**

#### Measuring Time by Interval Counting [间隔计数]

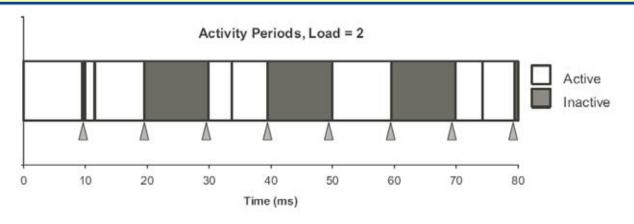
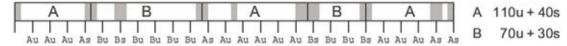


Figure 9.6: Graphical Representation of Activity Periods for Trace in Figure 9.5. Timer interrupts are indicated by gray triangles





#### (b) Actual Times

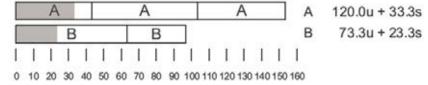


Figure 9.7: Process Timing by Interval Counting. With a timer interval of 10 ms, every 10 ms segment is assigned to a process as part of either its user (u) or system (s) time. This accounting provides only an approximate measure of program execution time.







### **Measurement and Profiling**

#### Measuring Time by Interval Counting [间隔计数]

# unix> time prog -n 17

2.230u 0.260s 0:06.52 38.1% 0+0k 0+0io 80pf+0w

- ☐ The first three numbers shown in this line are times.
- ☐ The first two show the seconds of user and system time.
- ☐ With a timer interval of 10 ms, all timings are multiples of hundredths of seconds. The third number is the total elapsed time.
- □ Observe that the system and user time sum to 2.49 seconds, less than half of the elapsed time of 6.52 seconds.
- ☐ The percentage indicates what fraction the combined user and system times were of the elapsed time, e.g. (2.23+0.26)/6.52=0.381
- ☐ The remaining statistics summarize the paging and I/O behavior.







Measuring Time by Interval Counting [间隔计数]

Do U remember the RECURSION example?

Gzhu\$ time ./recursion 10000

18%

0m0.456s real user 0m0.031s 0m0.052s SYS







# Timing Mechanisms:

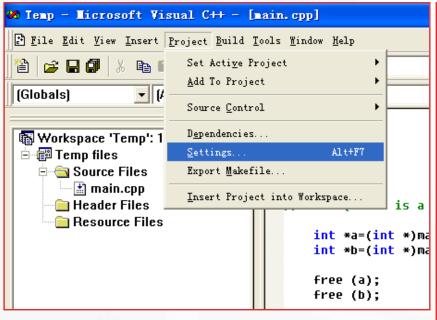
- Using timer directly
  - See precise.zip (downloadable).
- Using statistical sampling
  - Timer periodically interrupts the program and records the program counter or increments a counter representing a range of program counters
- Not necessary to implement from scratch.
  - **Profiling**







# Setting profiling:



| General   Debug   C/C++ Link   Resources   B   |
|--|
| Category: General ▼ Reset  |
| Output file <u>n</u> ame:  |
| Debug/Temp.exe   |
| Object/ <u>l</u> ibrary modules:   |
| kernel32.lib user32.lib gdi32.lib winspool.lib comdlg32.lib  |
| ✓ Generate debug info  |
| ☐ Link incrementally ☐ Generate <u>m</u> apfile  |
|  |
| Project Options:   |
| kernel32.lib user32.lib gdi32.lib winspool.lib<br>comdlg32.lib advapi32.lib shell32.lib ole32.lib<br>oleaut32.lib uuid.lib odbc32.lib odbccp32.lib |









## In Visual Studio .Net C++:

| Add New Project                                 | İ                       |  |                             |    | ?×     |
|---|-------------------------|--|-----------------------------|----|--------|
| Project types: Te  Visual C++ CLR Win32 General |                         | Templates:   |                             |    |        |
|   |                         | Visual Studio installed templates  Class Library CLR Empty Project  My Templates | pplication<br>s Application |    |        |
|   |                         | Search Online Templates  |                             |    |        |
| An empty project fo                             | or creating a local app | lication   |                             |    |        |
| Name:   | timer-managed           |  |                             |    |        |
| Location:                                       | C:\cygwin\home\ve       | erasam\ssd6\SSD6   |                             | ~  | Browse |
|   |                         |  |                             | ОК | Cancel |









### In Visual Studio .Net C++:

See: profiler-mod

- open a MS-DOS command prompt window, and run the EnableProfiler.bat
- 2. The batch file will register the profiler in your computer's registry.
- 3. Then, run your managed application (\*.exe ) from the same MS-DOS window to profile your program.
- 4. The profiling data will be written to a file named output.log.







### In Visual Studio .Net C++:

- ☐ Thread ID: The thread under which the function executed
- ☐ Function: Name of the function
- ☐ Times Called: The number of times the function was called
- ☐ Exclusive Time: Amount of time (in seconds) spent in the function excluding time spent in its callees, Suspended Time and Profiler Time
- ☐ Callee Exclusive Time: Amount of time (in seconds) spent in the function and its callees (children) excluding Suspended Time and **Profiler Time**
- ☐ Inclusive Time: Amount of time (in seconds) spent in the function including Callee Time, Suspended Time and Profiler Time
- ☐ Callee Time: Amount of time (in seconds) spent in the callees (including Suspended Time and Profiler Time spent under the callees)
- ☐ Suspended Time: Amount of suspended time (in seconds)
- ☐ Profiler Time: Amount of time spent by profiler (in seconds)







## **Measurement and Profiling**

# Profiling [summary]:

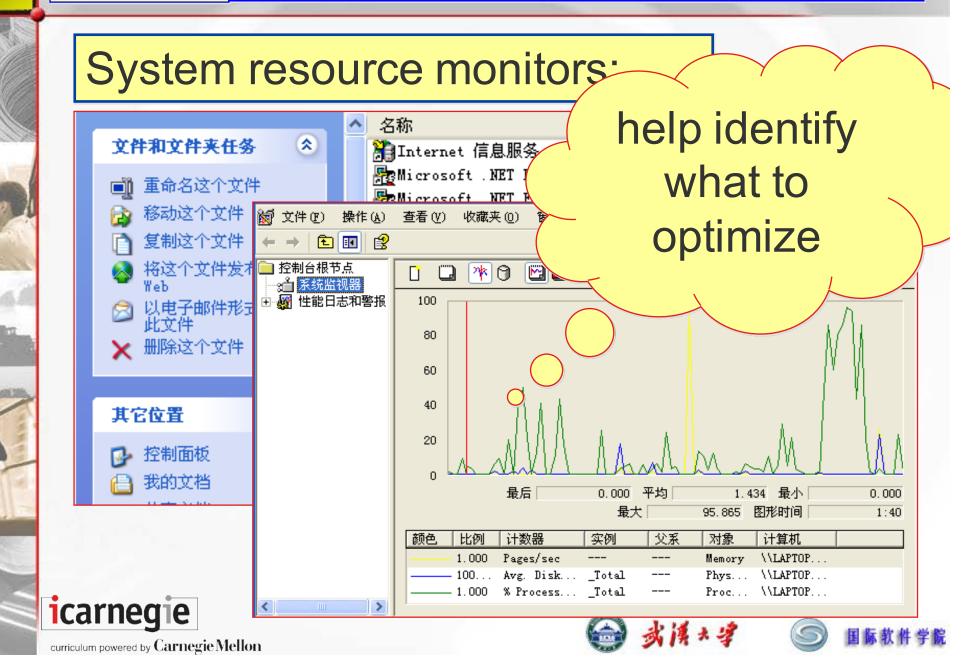
- Program performance can be measured
  - timing the execution
  - statistical sampling (profiling).
- Timers to measure the run time
- Profiling to help identify what to optimize
- System resource monitors:
  - CPU time
  - memory usage
  - I/O behavior







## **Measurement and Profiling**



# What's a Hot Spot?

- The 80/20 Rule
  - 80% of the CPU time is spent in 20% of the program
  - These places where the computer spends most of its time are also called *hot spots*, inner loops, and kernels.







# **Hot Spots**

- All programs have hot spots
- A little effort can have large payoffs.
- If only part of a program can be sped up with parallel processing, then the sequential part of the program would determine and limit the possible speed up.
- Knowing how to optimize is an important skill. Knowing when to stop optimizing is also important







# **An example for Hot Spots**

Next Time...







# Lecture 3 Next Time.....

- Read:
  - Unit 4 (Go through it)

- Do:
  - Multiple-Choice Quizzes in Unit 4







### Lecture 3 In the Lab.....



- Get debugging malloc.zip from ftp@iss
- Define your own Malloc() & Free()
- Define your own Structure for Payload

| Header         |         | Footer |  |  |
|----------------|---------|--------|--|--|
| Checksum Fence | Payload | Fence  |  |  |









## The Contents in SSD6 cover:

4.1 Measurement and Profiling
Where to start? Where to stop? How-to?

4.2 Hot Spots

Find the weak-knees!

4.3 Practical Hints

Take a tour...







# **Hot Spots**

- All programs have hot spots
- A little effort can have large payoffs.
- If only part of a program can be sped up with parallel processing, then the sequential part of the program would determine and limit the possible speed up.
- Knowing how to optimize is an important skill. Knowing when to stop optimizing is also important







$$T_{New} = (1 - \alpha)T_{Old} + (\alpha T_{Old})/k$$

$$S = \frac{T_{Old}}{T_{New}} = \frac{1}{(1-\alpha) + \alpha/k}$$







$$S = \frac{T_{Old}}{T_{New}} = \frac{1}{(1 - \alpha) + \alpha / k}$$

$$S_{\infty} = \frac{1}{(1-\alpha)}; \ k \to \infty$$





# An example for Hot Spots

**A Trigram Game** 

This is an example. The bird is an eagle. The boy is an expert.

The bird is an expert. This is an eagle. The boy is an example.







# **A Trigram Game**

See:

Words.cpp Words3.cpp







# **A Trigram Game**

|    | В   | С               | D                 | E                           | F              | G           | Н                     | I                | J      | K               |
|----|---|-----------------|-------------------|-----------------------------|----------------|-------------|-----------------------|------------------|--------|-----------------|
| 1  | Function  | Times<br>Called | Exclusive<br>Time | Callee<br>Exclusive<br>Time | Inclusive Time | Callee Time | Suspe<br>nded<br>Time | Profiler<br>Time | Func%  | Func<br>+child% |
| 2  | static void <module>::main( )</module>              | 1               | 0.004192          | 951.23658                   | 4732.46844     | 4732.46385  | 0                     | 0.000398         | 0.00%  | 100.00%         |
| 3  | static void <module>::generate( void, in</module>   | 1               | 0.384371          | 949.754105                  | 4725.739688    | 4725.355035 | 0                     | 0.000282         | 0.06%  | 99.84%          |
| 4  | static void <module>::countpairs( int32</module>    | 9998            | 421.458386        | 631.672108                  | 3152.368877    | 2730.597402 | 0                     | 0.313089         | 66.25% | 66.41%          |
| 5  | static void <module>::pick_nth( *, int8, *</module> | 9998            | 213.207295        | 315.888776                  | 1570.155352    | 1356.638362 | 0                     | 0.309695         | 33.51% | 33.21%          |
| 6  | static void <module>::readwords( )</module>         | 1               | 0.095483          | 1.476856                    | 6.722488       | 6.626706    | 0                     | 0.000298         | 0.02%  | 0.16%           |
| 7  | static void <module>::getword( int32 )</module>     | 8254            | 0.760527          | 1.127512                    | 5.373768       | 4.356743    | 0                     | 0.256499         | 0.12%  | 0.12%           |
| 8  | static void <module>::string_to_heap( *</module>    | 8253            | 0.161802          | 0.25386                     | 1.252938       | 0.838691    | 0                     | 0.252445         | 0.03%  | 0.03%           |
| 9  | void System.AppDomain::SetupDomain                  | 1               | 0.024919          | 0.090077                    | 0.098413       | 0.073145    | 0                     | 0.000349         | 0.00%  | 0.01%           |
| 10 | void System AppDomain: SetupFusionS                 | 1               | 0.056852          | 0.064624                    | 0.072301       | 0.013876    | n                     | 0.001574         | 0.01%  | 0.01%           |









# Lecture 6 Hot Spots

```
A Trigram Game
 char *string_to_heap(char *s)
    for (int i = 0; i < wordcount; i++) {
         if (streql(s, words[i])) {
           return words[i];
   char *h = (char *) malloc(strlen(s) + 1);
   if (!h) fatal("no more memory for words");
   strcpy(h, s);
   return h;
                    char *string_to_heap(char *s) {
                      int slen = strlen(s)+1;
                      char *h = (char *) malloc(slen);
                      if (!h) fatal("no more memory for words");
                      strcpy_s(h, slen, s);
                      return h;
icarnegie
```





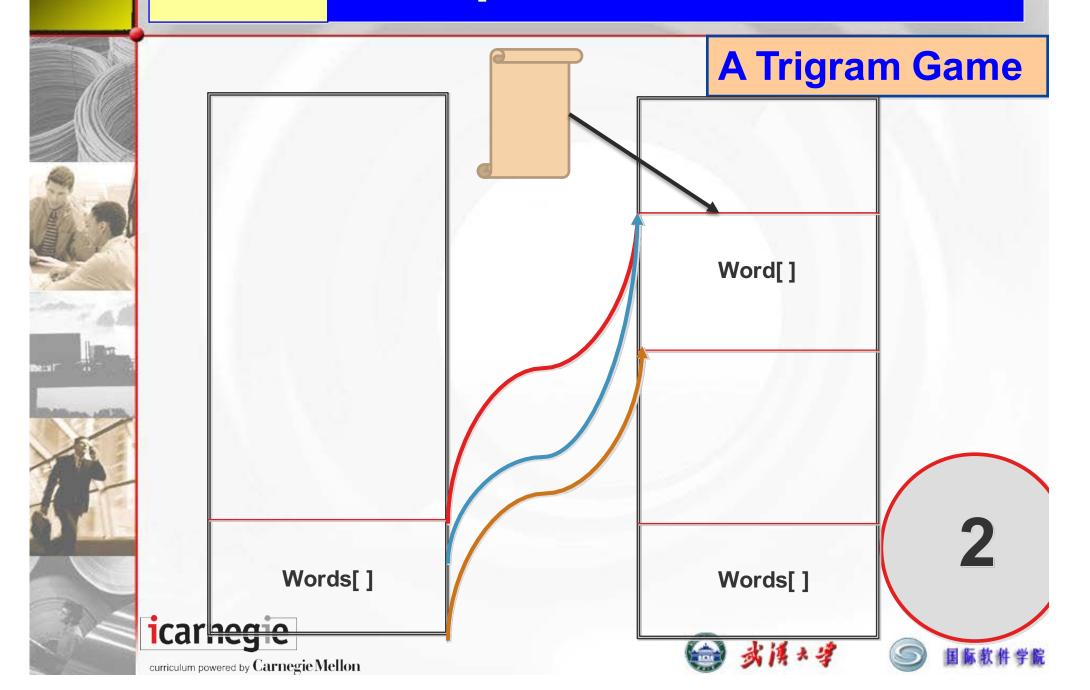


# Lecture 6 Hot Spots

```
#define streql(s1, s2) (strcmp(s1, s2) == 0)
                                           A Trigram Game
     int countpairs(char *word1, char *word2) {
       int count = 0; int i;
       for (i = 0; i < wordcount; i++) {
          // find word1.
         if (streql(words[i], word1)) {
           // IT Toung, check IT IT'S Totlowed by word2:
            if (streql(words[(i + 1) % wordcount], word2)) {
            // yes, we found (word1, word2) count++;
       return count;
                   int countpairs(char *word1, char *word2) {
                      int count = 0; int i;
                      for (i = 0; i < wordcount; i++) {
                        // find word1:
                         if (words[i] == word1) {
                          // IT Tound, check IT It's followed by wo
                          if (streql(words[(i + 1) % wordcount], wo
icarnegie
                          // yes, we found (wint) * word 2 count # 常能
curriculum powered by Carnegie Mellon
```

11 Lecture 6 Hot Spots **A Trigram Game** Word[] Words[] Words[] curriculum powered by Carnegie Mellon

# Lecture 6 Hot Spots



13

# Lecture 6 Hot Spots

|    | В   | С               | D                 | E                   | F             | G           | Н                     | I                | J      | K               |
|----|---|-----------------|-------------------|---------------------|---------------|-------------|-----------------------|------------------|--------|-----------------|
| 1  | Function  | Times<br>Called | Exclusive<br>Time | Callee<br>Exclusive | nclusive Time | Callee Time | Suspe<br>nded<br>Time | Profiler<br>Time | Func%  | Func<br>+child% |
| 2  | static void <module>::main( )</module>              | 1               | 0.004192          | 951.23658           | 4732.46844    | 4732.46385  | 0                     | 0.000398         | 0.00%  | 100.00%         |
| 3  | static void <module>::generate( void, in</module>   | 1               | 0.384371          |                     | 4725.739688   | 4725.355035 | 0                     | 0.000282         | 0.06%  | 99.84%          |
| 4  | static void <module>::countpairs( int32</module>    | 9998            | 421.458386        | 631.672108          | 3152.368877   | 2730.597402 | 0                     | 0.313089         | 66.25% | 66.41%          |
| 5  | static void <module>::pick_nth( *, int8, *</module> | 9998            | 213.207295        | 315.888776          | 1570.155352   | 1356.638362 | 0                     | 0.309695         | 33.51% | 33.21%          |
| 6  | static void <module>::readwords( )</module>         | 1               | 0.095483          | 1.476856            | 6.722488      | 6.626706    | 0                     | 0.000298         | 0.02%  | 0.16%           |
| 7  | static void <module>::getword( int32 )</module>     | 8254            | 0.760527          | 1.127512            | 5.373768      | 4.356743    | 0                     | 0.256499         | 0.12%  | 0.12%           |
| 8  | static void <module>::string_to_heap( *</module>    | 8253            | 0.161802          | 0.25386             | 1.252938      | 0.838691    | 0                     | 0.252445         | 0.03%  | 0.03%           |
| 9  | void System.AppDomain::SetupDomain                  | 1               | 0.024919          | 0.090077            | 0.098413      | 0.073145    | 0                     | 0.000349         | 0.00%  | 0.01%           |
| 10 | void System AppDomain: SetupFusionS                 | 1               | 0.056852          | 0.064624            | 0.072301      | 0.013876    | n                     | 0.001574         | 0.01%  | 0.01%           |

| 1  | Function   | Times<br>Called | Exclusive<br>Time | Callee<br>Exclusive | Inclusive<br>Time | Callee Time | Suspe<br>nded<br>Time | Profiler<br>Time | Func%  | Func<br>+child% |
|----|--|-----------------|-------------------|---------------------|-------------------|-------------|-----------------------|------------------|--------|-----------------|
| 2  | static void <module>::main( )</module>           | 1               | 0.003923          | 89.564498           | 411.379271        | 411.374964  | 0                     | 0.000385         | 0.01%  | 100.00%         |
| 3  | static void <module>::readwords( )</module>      | 1               | 0.093546          |                     | 407.500631        | 407.406796  | 0                     | 0.000289         | 0.15%  | 97.47%          |
| 4  | static void <module>::string_to_heap(</module>   | 8253            | 60.626506         | 86.131939           | 402.416761        | 341.549836  | 0                     | 0.240419         | 96.51% | 96.17%          |
| 5  | static void <module>::generate( void,</module>   | 1               | 0.296708          | 2.261382            | 3.872616          | 3.575621    | 0                     | 0.000287         | 0.47%  | 2.52%           |
| 6  | static void <module>::getword( int32 )</module>  | 8254            | 0.700869          | 1.072232            | 4.990035          | 4.049166    | 0                     | 0.24             | 1.12%  | 1.20%           |
| 7  | static void <module>::countpairs( int3</module>  | 9998            | 0.751609          | 0.751609            | 1.068417          | 0           | 0                     | 0.316809         | 1.20%  | 0.84%           |
| 8  | static void <module>::pick_nth( *, int8</module> | 9998            | 0.317123          | 0.317123            | 0.623271          | 0           | 0                     | 0.306148         | 0.50%  | 0.35%           |
| 9  | void System.AppDomain::SetupDomai                | 1               | 0.004213          | 0.030956            | 0.039072          | 0.034516    | 0                     | 0.000343         | 0.01%  | 0.03%           |
| 10 | void System.AppDomain::SetupFusion               | 1               | 0.018392          | 0.026332            | 0.033811          | 0.01384     | 0                     | 0.001579         | 0.03%  | 0.03%           |
| 11 | Livia Contra Ambridation Cation                  |                 | 0.000002          | 0.000070            | 0.007004          | 0.0004.0    | _ ^                   | 0.004.000        | 0.000  | 0.000           |



icarnegie After optimized!







- Some Expensive Operations in C
  - printf()
  - malloc()
  - new()









- *Printf()* causes *malloc* to be called in the usual implementation, adding an unexpected cost
- String manipulation is expensive, whether it is formatting text as in *printf*, reading ASCII text and converting to numbers, or performing string comparisons
- Modern computers are designed to access memory a word at a time. Accessing a character is no faster than accessing a word, and may be slower
- String operations usually operate one character at a time, i.e strlen(), has to examine every character, one at a time









- use static or local variables to avoid allocating memory on the heap;
- keep a list of objects that need to be allocated often → explicit free-list
- Then, allocation is just a matter of removing an object from the list, and freeing simply inserts the object on the list







### **Lecture 6**

# **Practical Hints**



```
void main()
   cout << "Allocated: ";
   for (int i = 0; i < 10; i++) {
        a[i] = Myobject::alloc();
      cout << a[i] << " ";
    cout << endl;
    // free the even-numbered objects:
    cout << "Freed: ";
    for (i = 0; i < 5; i++) {
         cout << a[i * 2] << " ";
         a[i * 2]->free();
    cout << endl;
    // reallocate the even-numbered objects:
    for (i = 0; i < 5; i
       a[i * 2] = Myobject::alloc();
     cout << a[i * 2] << " ",
    cout << endl;
```







### **Lecture 6**

# **Practical Hints**

```
iostream>
using namespace std;
Myobject{
public:
   Myobject*alloc();
    void free();
    Myobject*next; // used for allocation
    Myobject*freelist;
   long x; // your member variables go here,
"declare and initialize empty free list:
Myobject*Myobject::freelist = NULL;
```





### **Lecture 6**

curriculum powered by Carnegie Me

# **Practical Hints**

```
icarnegie
```

```
Myobject*Myobject::alloc()
     Myobject*p;
     // fast allocation: pop object from freelist
      if (freelist)
           p = freelist;
           e {

Myobject;
      } else {
       return p;
  void Myobject::free()
       this- freelist;
        freelist = this;
```

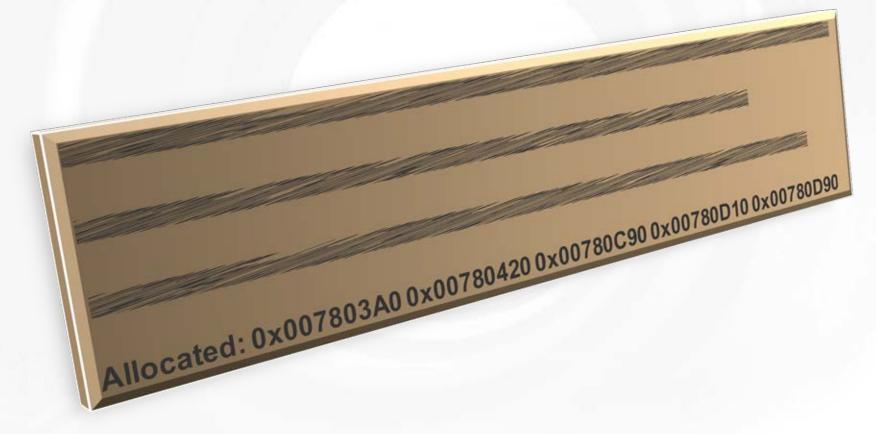








Output:









# Lecture 6 Practical Hints {more}

## **Fast Allocation and Free**

Use a large block of memory from which smaller chunks are allocated to compute some result. After the result is obtained, the entire block is freed. This is fast because:

- Memory is allocated simply by incrementing the "free" pointer by the number of bytes you need to allocate.
- There is no need to free each allocated object; you free all objects at once by freeing the entire pool at once.

An Allocator!







## Lecture 6 About Excise 4 ......



- Get substitution.c and data.zip from ftp@iss
- Goto support.microsoft.com, search for an article with Article ID Q224382
- Optimize code to get better performance
- Measure performance by profiling





